

Experiment No. 2: Deep neural network on IMDB Dataset

1. The dataset is the Large Movie Review Dataset, often referred to as the IMDB dataset.
2. The IMDB dataset contains 50,000 highly polar movie reviews (good or bad) for training and the same amount again for testing. The problem is to determine whether a given movie review has a positive or negative sentiment.

#Load the IMDB Dataset with Keras

```
import numpy as np

from tensorflow.keras.datasets import imdb

import matplotlib.pyplot as plt

# load the dataset

(X_train, y_train), (X_test, y_test) = imdb.load_data()

X= np.concatenate((X_train, X_test), axis=0)

y= np.concatenate((y_train, y_test), axis=0)

# summarize size

print("Training data: ")

print(X.shape)

print(y.shape)
```

OUT:

Training data:

```
(50000,)
```

```
(50000,)
```

```
#Word Embedding
```

```
imdbload_data (nb_words=5000)
```

```
#truncate or pad the dataset to a length of 500 for each observation
```

```
X_train = sequence.pad_sequences(X_train, maxlen=500)
```

```
X_test = sequence.pad_sequences(X_test, maxlen=500)
```

```
Embedding(5000, 32, input_length=500)
```

```
#Simple Multi-Layer Perceptron Model for the IMDB Dataset
```

```
# MLP for the IMDB problem
```

```
from tensorflow.keras.datasets import imdb
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense
```

```
from tensorflow.keras.layers import Flatten
```

```
from tensorflow.keras.layers import Embedding
```

```
from tensorflow.keras.preprocessing import sequence
```

```
# load the dataset but only keep the top n words, zero the rest
```

```
top_words = 5000
```

```
(X_train, y_train), (X_test, y_test) = imdb.load_data (num_words=top_words)
```

```
#bound reviews at 500 words, truncating longer reviews and zero-padding shorter one
```

```
max_words =500
```

```
X_train = sequence.pad_sequences (X_train, maxlen=max_words)
```

```
X_test = sequence.pad_sequences (X_test, maxlen=max_words)
```

```
# create the model
```

```
model = Sequential()
```

```
model.add (Embedding (top_words, 32, input length=max_words))
```

```
model.add (Flatten ())
```

```
model.add (Dense [250, activation='relu'])
```

```
model.add (Dense(1, activation='sigmoid'))
```

```
model.compile (loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
model.summary()
```

```
# Fit the model
```

```
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=2, batch_size=128,  
verbose=2)
```

```
# Final evaluation of the model
```

```
Scores = model.evaluate(X_test, y_test, verbose=0)
```

```
print("Accuracy: %.2f%%" % (scores[1]*100))
```

OUT:

Epoch ½

196/196-4s - loss: 0.5579 - accuracy: 0.6664 - val loss: 0.3060 - val_accuracy: 0.8700 -

4s/epoch - 20ms/step

Epoch 2/2

196/196 - 4s - loss: 0.2108 - accuracy: 0.9165 - val_loss: 0.3006 - val_accuracy: 0.8731 -

4s/epoch - 19ms/step

Accuracy: 87.319%