

Experiment No. 4 : Classification of MNIST Fashion dataset using CNN

```
import matplotlib.pyplot as plt

import numpy as np

import pandas as pd

from keras.utils import to_categorical

from matplotlib.pyplot import figure, show

import warnings

import seaborn as sns

warnings.filterwarnings(ignore')

import matplotlib.style as style

from sklearn.model_selection import train_test_split

from keras.layers import Input, Concatenate, concatenate, Dense, Embedding, Dropout,
Conv2D, MaxPooling2D

from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau, History

from keras.layers import Dropout, Flatten, GlobalAveragePooling2D, Activation

from sklearn.preprocessing import LabelEncoder, OneHotEncoder

from keras.preprocessing.image import imageDataGenerator

from sklearn.model_selection import train_test_split

from keras.applications.resnet50 import ResNet50

from keras.callbacks import ReduceLROnPlateau

from keras.callbacks import ModelCheckpoint

from keras.applications.vgg16 import VGG16

from keras.utils import to_categorical

from sklearn.utils import class_weight

from keras.layers.normalization import BatchNormalization

from matplotlib import pyplot as plt
```

```

from keras import backend as K

from keras.optimizers import SGD

from keras.models import Model

import seaborn as sns

import numpy as np

import argparse

import time

import glob

import cv2

import numpy

import os

import glob

import sys

import os

import json

import pprint

import warnings

warnings.filterwarnings('ignore')


#Load Data

!curl -L -O https://www.dropbox.com/s/heyql2my8uwotq/fashionmaistzip

!unzip fashionmnist.zip


#Load training and test data using dataframes from Pandas.

train = pd.read_csv("fashion-mnist_train.csv")

test = pd.read_csv("fashion-mnist_test.csv")

```

```
Img_rows, img_cols = 28, 28
```

```
input_shape = (img_rows, img_cols, 1)
```

```
X= train.iloc[:,1:]
```

```
Y= train.iloc[:,1]
```

```
X_test = test.iloc[:, 1:]
```

```
Y_test = test.iloc[:, 1]
```

```
#Normalization
```

```
X= np.asarray(X).reshape (X.shape [0], img_rows,img_cols, 1)
```

```
X_test = np.asarray(X_test).reshape(X_test.shape [0], img_rows,img_cols,1)
```

```
X= (255. - X) /255.
```

```
X_test = (255. - X_test) / 255.
```

```
#Number of classes
```

```
classes = len(Y['label'],value _counts())
```

```
print("Number of features: ", X.shape[1])
```

```
print("Number of train samples: ", Xshape [0])
```

```
print("Number of test samples: ", X test.shape [0])
```

```
OUT:
```

```
Number of features: 28
```

```
Number of train samples: 60000
```

```
Number of test samples: 10000
```

```
#Training
```

```

Y_test = to_categorical(Y_test)

Y= to_categorical (Y)

X_train, X_val, Y_train, Y_val = train_test_split(X, Y, stratify=Y, test_size=0.2,
random_state=66)

Irr = ReduceLROn Plateau (monitor='val_loss', factor=0.1, patience=2, verbose=1,
epsilon=1e-3, mode= 'min')

early_stopping = EarlyStopping(monitor='val loss',patience=5,verbose=0, mode='auto')

checkpoint = ModelCheckpoint("checkpoint.hdf5", monitor='val_acc', verbose=1,
save_best_only=True, mode='max')

batch size = 64

epochs = 10

```

```

from sklearn.model_selection import GridSearchCV

from keras.wrappers.scikit_learn import KerasClassifier

```

```

# define the grid search parameters

batch_size = [16, 32, 64, 80]

epochs = [10, 25, 50]

param_grid = dict (batch_size=batch_size, epochs=epochs)

model = KerasClassifier(build_fn=model_basic, verbose=0)

Grid=GridSearchCV(estimator=model, param_grid=param_grid,n_jobs1, cv=3)

grid_result = grid.fit(X_train, Y_train)

```

```

# summarize results

print("Best: %using %s" % (grid_result.best_score, grid_result.bestparams_))

means = grid_result.cv_results ['mean_test_score']

```

```

stds = grid_result.cv_results ['std_test_score']

params =grid_result.cv_results ['params']

for mean, stdev, param in zip(means, stds, params):

print("%f (%f) with: %r" % (mean, stdev, param))


#Model

def model_basic(classes=classes,optimizer='adam'):

kernel_size = (3,3)

dropout = 0.25

pool_size = (2,2)

inputs = Input(shape=(img_rows, img_cols, 1))

y= Conv2D (filters=32, kernel_size=kernel_size,activation='relu',padding='same') (inputs)

y= MaxPooling2D(pool_size=pool_size,strides=(2,2)) (y)

y= Dropout(dropout)(y)


y= Flatten()(y)


y= Dense(256,activation='relu')(y)

y= BatchNormalization()(y)

y= Dropout(dropout)(y)

outputs = Dense(classes, activation='softmax')(y)


model = Model(inputs=inputs, outputs=outputs)

model.compile (optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

return model

basic_model = model_basic()

```

```
import warnings

warnings.filterwarnings('ignore')

history = basic_model.fit(X_train, Y_train, batch_size=batch_size, epochs=epochs,
verbose=1, validation_data=(X_val, Y_val))
```

OUT:

Train on 48000 samples, validate on 12000 samples

Epoch 1/10

48000/48000 [=====]-82s 2ms/step- loss: 0.4390 - acc:.8470 - val_loss: 0.4019val _acc:
0.8618

Epoch 2/10

48000/48000 [=====]-82s 2ms/step - loss: 0.3458 - acc: 0.8776 - val_loss: 0.3046
val_acc: 0.8932

Epoch 3/10

48000/48000 [=====] -82s 2ms/step - loss: 0.3110 - acc: 0.8883 - val_loss: 0.2947 -
val_acc: 0.8953

Epoch 4/10

48000/48000 [====]-81s 2ms/step - loss: 0.2935 - acc: 0.8937 - val_loss: 0.2772 -
val_acc: 0.9024

Epoch 5/10

48000/48000 [=====]-85s 2ms/step- loss: 0.2710 - acc: 0.9030 - val_loss: 0.2855-
val_acc: 0.8952

Epoch 6/10[=====| -84s 2ms/step - loss: 0.2592 - acc: 0.9067 - val _loss: 0.2574-
val_acc: 0.9063

Epoch 7/10

48000/48000 [=====]-85s 2ms/step-loss: 0.2450 - acc: 0.9107 - val_loss: 0.2773 -

val acc: 0.8998

Epoch 8/10

48000/48000 [=====] - 84s 2ms/step - loss: 0.2338 - acc: 0.9147 - val_loss: 0.2833 -

val_acc: 0.8955

Epoch 9/10

48000/48000 [=====]-82s 2ms/step - loss: 0.2239 - acc: 0.9174 - val_loss: 0.2553 -

val_acc: 0.9127

Epoch 10/10

48000/48000 [=====]-84s 2ms/step - loss: 0.2153 - acc: 0.9207 - val loss: 0.2564 -

val_acc: 0.9123

```
score = basic_model.evaluate(X_test,Y_test, verbose=0)
```

```
print("Test loss:", score[0])
```

```
print("Test accuracy:", score[1])
```

OUT:

Test loss: 0.2463475521683693

Test accuracy: 0.9137